

## Set 4 - MPI Part II

Issued: November 23, 2022

Hand in (optional): December 6, 2022 08:00

**Grading:** To get full credits solve any two of the questions.

### Question 1: Diffusion (100 points)

The diffusion of a substance can be described by the equation

$$\frac{\partial c(x, y, t)}{\partial t} = D \left( \frac{\partial^2 c(x, y, t)}{\partial x^2} + \frac{\partial^2 c(x, y, t)}{\partial y^2} \right),$$

where  $c$  is the concentration of the substance at position  $(x, y)$  and at time  $t$ , and  $D$  is the diffusion constant. The diffusion process happens in the domain  $|x| < L/2$  and  $|y| < L/2$ . The concentration is zero on the boundaries of the domain. The initial concentration is

$$c(x, y, 0) = \begin{cases} 1, & \text{if } |x| < L/4 \text{ and } |y| < L/4, \\ 0, & \text{otherwise.} \end{cases}$$

- a) The skeleton code solves the equation on a uniform grid using a central finite difference scheme in space and forward Euler time integration. Parallelize the code by filling parts marked by `TODO` in the functions `advance` and `main`. Use a tiling decomposition scheme (i.e., distribute the rows evenly to the MPI processes). How to run the code:
- `make` to compile the code
  - `make run` to run single core (please not in the login node on euler!). If not locally on the laptop, use `sbatch launch_single.sh` to submit a job via slurm.
  - to run multicore use: `mpirun -n x ./diffusion D L N`. You can also use `sbatch launch_single.sh` to submit a job via slurm.
- b) For a given time compute the integral of  $c(x, y, t)$  over the domain (total amount of the substance). Fill the missing MPI parts in `compute_diagnostics`, and plot the result as a function of time using  $D = 1$ ,  $L = 2$  and  $N = 100$ . When run correctly, the code will output file called `diagnostics.dat`. To plot this data use `python plot_diagnostics.py` (module load python).
- c) For a given time compute the histogram of  $c(x, y, t)$  in the function `compute_histogram` by implementing the missing MPI parts marked by `TODO`, and plot or print the resulting histogram for  $t = 0.5$  using  $D = 1$ ,  $L = 2$  and  $N = 100$ .

- d) Suggest other ways to divide the real-space domain between processes with the aim of minimizing communication overhead. Prove your argument by computing the message communication size for the tiling domain decomposition and for your suggestion.
- e) Make a strong and weak scaling plot up to 48 cores. Justify what is happening in your plots. Make at least five different numbers of cores runs (e.g. [1, 12, 24, 36, 48] or [1, 2, 4, 8, 16]). For the strong scaling plot use:  $N = \{1024, 2048, 4096, 8192\}$ , in other words, plot at least four lines (if too slow use smaller  $N$ 's). For the weak scaling plot, use  $N = 1024$  and  $N = 2048$  (if those are taking too long use smaller  $N$ 's). Use  $D = 1$ ,  $L = 2$  and modify the number of timesteps  $\text{step} = 100$ . Do not forget to state which CPU you ran the tests on! You can use the `run.sh` script to run for different number of cores. On euler use `run.sh` via `sbatch launch.sh`. Draw the plots either directly on the paper (which is the way you will do it on the exam). Or use any desired plotting scripts.