

**Set 6 – Sparse Linear Algebra with MPI**

Issued: December 07, 2022

Hand in (optional): December 20, 2022 23:59

**Question 1: CSR format (10 points)**

Sparse matrices have a limited number of nonzero elements. Specialized formats are used for efficient storage and operations on them. The CSR format (compressed sparse row) represents a matrix  $A_{ij}$  by three one-dimensional arrays:

- $A[k]$ , nonzero coefficients in the row-major order,  $0 \leq k < N_z$ , where  $N_z$  is the number of non-zero elements of the matrix;
- $K[i]$ , extents of rows, row  $i$  consists of  $A[K[i]:K[i+1]]$ ;
- $J[k]$ , column indices.

Write down a representation of the following matrix in the CSR format

$$A = \begin{bmatrix} 2 & -1 & 0 & -1 \\ -1 & 2 & -1 & 0 \\ 0 & -2 & 4 & -2 \\ -1 & 0 & -1 & 2 \end{bmatrix}.$$

From that representation, compute the product  $Au$  with column  $u = [0, 1, 2, 3]^T$ .

**Question 2: Matrix-vector product with MPI (90 points)**

Matrices and vectors can be distributed to multiple processors. Consider the matrix-vector product

$$b = Au$$

with a square matrix  $A \in \mathbf{R}^{n \times n}$  and vectors (columns)  $u, b \in \mathbf{R}^n$ . Each of  $p$  processors stores  $n/p$  rows of the matrix, the corresponding chunks of vector  $b$  and the same chunks of vector  $u$ . For dense matrices (many non-zero elements), the communication pattern is all-to-all. However, for a banded matrix (e.g. tridiagonal) communication is required only between neighbouring chunks. Such sparse matrices (most elements are zero) result from the discretization of Partial Differential Equations (PDEs).

One example is the two dimensional diffusion equation

$$\frac{\partial u}{\partial t} = \frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2}$$

which can be discretized as

$$u_{i,j}^{n+1} = u_{i,j}^n + \frac{\Delta t}{h} (u_{i-1,j}^n + u_{i+1,j}^n + u_{i,j-1}^n + u_{i,j+1}^n - 4u_{i,j}^n)$$

for  $i = 0, \dots, N_i - 1$ ,  $j = 0, \dots, N_j$ . This can be rewritten as a matrix-vector product  $u^{n+1} = u^n + Au^n$ , by defining

$$u^n = [u_{0,0}^n, u_{0,1}^n, \dots, u_{0,N_j-1}^n, u_{1,0}^n, u_{1,1}^n, \dots, u_{1,N_j-1}^n, \dots, u_{N_i-1,0}^n, u_{N_i-1,1}^n, \dots, u_{N_i-1,N_j-1}^n]^T.$$

- a) Derive matrix  $A$  so that the matrix-vector product  $u^{n+1} = u^n + Au^n$  is equivalent to the discrete diffusion equation above. Matrix  $A$  is known as the Poisson matrix. Assume periodic boundary conditions by *defining*

$$\begin{aligned} u_{N_i,j} &:= u_{0,j} & , j = 0, \dots, N_j - 1 \\ u_{-1,j} &:= u_{N_i-1,j} & , j = 0, \dots, N_j - 1 \\ u_{i,N_j} &:= u_{i,0} & , i = 0, \dots, N_i - 1 \\ u_{i,-1} &:= u_{i,N_j-1} & , i = 0, \dots, N_i - 1 \end{aligned}$$

- b) You are given a skeleton code that solves the diffusion equation by treating it as a series of matrix-vector products. Parallelize the sparse matrix-vector product with MPI. See the skeleton code for more details.

*Hint:* The algorithm should involve the following stages:

- traverse rows of the matrix, multiply the elements stored locally and collect global indices of columns that require communication; use `GlbToLoc()` to convert the global indices;
- use `GlbToRank()` to find the ranks of processors storing the required columns;
- use `MPI_Allreduce()` to compute the number of messages that every processor needs to receive;
- send and receive the indices of columns from other processors;
- send and receive the corresponding elements of vector  $u$ ;
- add the received data to the product.

- c) Report the weak scaling of your implementation on Euler.